

Dataset 4: Linear Discriminant Analysis

Heart Disease Risk Classification Study | n = 320 patients

1. Research Question

Research Question: Can a linear combination of seven clinical and lifestyle variables — Age, BMI, Blood Pressure, Cholesterol, Exercise Hours, Stress Level, and Smoking Years — effectively discriminate between patients at High Risk vs. Low Risk of heart disease? Which variables contribute most to this discrimination, and how does LDA compare to binary logistic regression for this classification task?

2. Variables

Role	Variable	High Risk Mean	Low Risk Mean	Difference
Outcome (DV)	Heart_Disease_Group	—	—	High Risk / Low Risk (binary)
Predictor	Age	53.61 yr	49.28 yr	+4.33 yr
Predictor	BMI	28.37	26.26	+2.11
Predictor	Blood_Pressure	133.12 mmHg	124.87 mmHg	+8.25 mmHg
Predictor	Cholesterol	211.05 mg/dL	199.56 mg/dL	+11.49 mg/dL
Predictor	Exercise_Hours_Per_Week	3.14 hr/wk	3.98 hr/wk	−0.84 hr/wk
Predictor	Stress_Level	5.47 (1–10)	4.59 (1–10)	+0.88
Predictor	Smoking_Years	8.72 yr	7.12 yr	+1.60 yr

Class balance: exactly 160 High Risk and 160 Low Risk patients (50/50) — perfectly balanced. High Risk patients are older, heavier, have higher blood pressure and cholesterol, exercise less, are more stressed, and have longer smoking histories.

3. Methodology — Linear Discriminant Analysis (LDA)

LDA finds the single linear combination of all 7 predictors that maximally separates the two risk groups. It does this by maximizing the ratio of between-group variance to within-group variance — producing the discriminant function D. All predictors are first standardized (z-scored) so coefficients are comparable. Data are split 75/25 (train/test, stratified). Performance is evaluated by accuracy, AUC-ROC, sensitivity, specificity, and F1 score. 10-fold stratified cross-validation tests generalizability. Both standardized (for ranking importance) and raw (for computing D from real patient values) coefficients are reported. Binary logistic regression is run on the same split for direct comparison.

4. Assumption Checks

Test	Result	Verdict
Levene — Age	F=0.273, p=0.602	✓ Equal variances
Levene — BMI	F=0.305, p=0.581	✓ Equal variances
Levene — Blood_Pressure	F=0.588, p=0.444	✓ Equal variances

Levene — Cholesterol	F=0.075, p=0.785	✓ Equal variances
Levene — Exercise_Hours_Per_Week	F=0.278, p=0.598	✓ Equal variances
Levene — Stress_Level	F=0.257, p=0.612	✓ Equal variances
Levene — Smoking_Years	F=6.358, p=0.012	✗ Minor violation — robust at n=320
Normality violations (Shapiro-Wilk) BMI, Stress, Smoking non-normal in some groups Minor — LDA robust at n=160/group		
Sample size: n=320, n/p = 45.7	Rule of thumb: n/p ≥ 5	✓ Adequate
Class balance: 160 / 160	50/50 split	✓ Perfect for LDA
Multicollinearity: all pairwise r < 0.44	No pair exceeds r=0.50	✓ No problematic collinearity

5. Standardized Discriminant Function Coefficients

Computed on z-scored predictors. Use to rank the RELATIVE IMPORTANCE of each variable. Larger |coefficient| = stronger discriminator. Sign indicates which direction increases the High Risk discriminant score.

Rank	Variable	Std Coefficient	Coef	Direction → High Risk
1	Exercise_Hours_Per_Week	+0.9567	0.9567	Less exercise → High Risk
2	Stress_Level	−0.9320	0.9320	Higher stress → High Risk
3	Blood_Pressure	−0.8570	0.8570	Higher BP → High Risk
4	Cholesterol	−0.8525	0.8525	Higher cholesterol → High Risk
5	BMI	−0.7689	0.7689	Higher BMI → High Risk
6	Age	−0.6091	0.6091	Older age → High Risk
7	Smoking_Years	−0.3438	0.3438	More smoking → High Risk

6. Raw (Unstandardized) Discriminant Function Coefficients

Computed in original measurement units. Use to calculate D for a real patient by plugging in their actual clinical values. D changes by [coefficient] for each one-unit increase in that variable, holding all others constant.

Variable	Raw Coefficient	Interpretation (per 1 unit increase)
Intercept (constant)	+20.809843	Baseline value when all predictors = 0
Age	−0.064963	D decreases 0.065 per year older
BMI	−0.157286	D decreases 0.157 per BMI unit
Blood_Pressure	−0.058819	D decreases 0.059 per mmHg
Cholesterol	−0.024090	D decreases 0.024 per mg/dL
Exercise_Hours_Per_Week	+0.647354	D increases 0.647 per additional hour/week of exercise
Stress_Level	−0.499347	D decreases 0.499 per stress unit (1–10 scale)
Smoking_Years	−0.061598	D decreases 0.062 per additional year of smoking

7. Discriminant Function Formulas

FORMULA A — Standardized (use to compare variable importance; input = z-scored):

$$D = (\approx 0.0000) + (+0.9567) \times z_{\text{Exercise_Hours_Per_Week}} + (-0.9320) \times z_{\text{Stress_Level}} + (-0.8570) \times z_{\text{Blood_Pressure}} + (-0.8525) \times z_{\text{Cholesterol}} + (-0.7689) \times z_{\text{BMI}} + (-0.6091) \times z_{\text{Age}} + (-0.3438) \times z_{\text{Smoking_Years}}$$

Where $z_X = (\text{Patient value} - \text{Sample mean}) / \text{Sample SD}$. The intercept ≈ 0 because the standardized function passes through the origin when priors are equal (50/50).

z-Score Reference Table (to use Formula A):

Variable	Sample Mean	Sample SD	z-formula
Age	51.45	9.40	$z = (\text{Age} - 51.45) / 9.40$
BMI	27.32	4.89	$z = (\text{BMI} - 27.32) / 4.89$
Blood_Pressure	129.00	14.58	$z = (\text{BP} - 129.00) / 14.58$
Cholesterol	205.31	35.41	$z = (\text{Chol} - 205.31) / 35.41$
Exercise_Hours_Per_Week	3.56	1.48	$z = (\text{Exer} - 3.56) / 1.48$
Stress_Level	5.03	1.87	$z = (\text{Stress} - 5.03) / 1.87$
Smoking_Years	7.92	5.58	$z = (\text{Smoke} - 7.92) / 5.58$

FORMULA B — Raw / Unstandardized (use to classify a real patient; plug in actual values):

$$D = (+20.809843) + (-0.064963) \times \text{Age} + (-0.157286) \times \text{BMI} + (-0.058819) \times \text{Blood_Pressure} + (-0.024090) \times \text{Cholesterol} + (+0.647354) \times \text{Exercise_Hours_Per_Week} + (-0.499347) \times \text{Stress_Level} + (-0.061598) \times \text{Smoking_Years}$$

Decision Rule (applies to both formulas):

D value	Classification	Meaning
$D < 0$	HIGH RISK ■	Patient sits on High Risk side of decision boundary
$D = 0$	Boundary	Maximum uncertainty — equal probability of either class
$D > 0$	LOW RISK ✓	Patient sits on Low Risk side of decision boundary

Group centroids on discriminant axis: High Risk = -0.7901 (negative side), Low Risk = 0.7177 (positive side). Separation = 1.5078 standardized units — the two groups are clearly on opposite sides of the $D=0$ boundary.

Worked Example — Hypothetical Patient:

Patient: Age=55, BMI=29.0, BP=135 mmHg, Cholesterol=215 mg/dL, Exercise=3.0 hrs/week, Stress=6.0, Smoking=9 years.

Step	Calculation	Contribution to D
Intercept		+20.8098
Age	$(-0.064963) \times 55$	-3.5730
BMI	$(-0.157286) \times 29.0$	-4.5613
Blood_Pressure	$(-0.058819) \times 135$	-7.9406
Cholesterol	$(-0.024090) \times 215$	-5.1794
Exercise	$(+0.647354) \times 3.0$	+1.9421
Stress_Level	$(-0.499347) \times 6.0$	-2.9961

Smoking_Years	$(-0.061598) \times 9$	-0.5544
TOTAL D		-2.0528 → HIGH RISK ■

8. Classification Results — Test Set (n = 80)

	Predicted HIGH RISK	Predicted LOW RISK	Total
Actual HIGH RISK	33 (True Positive)	7 (False Negative)	40
Actual LOW RISK	12 (False Positive)	28 (True Negative)	40
Total	45	35	80

Metric	Formula	Value	Interpretation
Accuracy	$(TP+TN)/(TP+TN+FP+FN) = (33+28)/80$	76.25%	Correctly classifies 3 in 4 patients
Sensitivity	$TP/(TP+FN) = 33/40$	82.50%	% of High Risk patients correctly caught
Specificity	$TN/(TN+FP) = 28/40$	70.00%	% of Low Risk patients correctly cleared
Precision	$TP/(TP+FP) = 33/45$	73.33%	% of High Risk predictions that were correct
F1 Score	$2 \times (Prec \times Sens) / (Prec + Sens)$	0.7765	Harmonic mean of precision & sensitivity
AUC-ROC	Area under the ROC curve	0.8350	83.5% chance of correct group ranking
CV Accuracy	10-fold cross-validation	75.00%±7.78%	Stable generalization — confirms the model

Clinical interpretation of false negatives: The 7 false negatives (FN) are High Risk patients classified as Low Risk — the most dangerous error in a clinical screening context. These patients receive no warning about their elevated cardiac risk. The model's sensitivity of 82.5% means 82.5% of High Risk patients are correctly flagged. To reduce false negatives further, the decision threshold could be lowered from D=0 to a negative value (e.g., D=−0.3), which would increase sensitivity at some cost to specificity.

9. Structure Matrix — Predictor–Discriminant Correlations

The structure matrix shows the Pearson correlation of each predictor with the discriminant score D. This is more stable than raw coefficients when predictors share moderate correlations, and provides the most reliable ranking of each variable's independent contribution to group separation. Threshold: $|r| \geq 0.30$ is considered practically significant.

Rank	Variable	r	p-value	Strength	Interpretation
1	Exercise_Hours_Per_Week	+0.4726	< 0.001	Strong	Less exercise → High Risk
2	Stress_Level	-0.4485	< 0.001	Strong	Higher → High Risk
3	Blood_Pressure	-0.3973	< 0.001	Moderate	Higher → High Risk
4	BMI	-0.3933	< 0.001	Moderate	Higher → High Risk
5	Age	-0.3526	< 0.001	Moderate	Higher → High Risk
6	Cholesterol	-0.2851	< 0.001	Moderate	Higher → High Risk
7	Smoking_Years	-0.1700	0.002	Weak	Higher → High Risk

10. LDA vs. Binary Logistic Regression — Full Comparison


Both methods were trained and tested on the identical 75/25 stratified split and evaluated with 10-fold cross-validation on the full dataset.

Metric	LDA	Logistic Reg	Difference	Winner
Accuracy (Test Set)	76.25%	76.25%	+0.00%	Tie
AUC-ROC	0.8350	0.8337	+0.0013	LDA ✓
Sensitivity (High Risk)	82.50%	82.50%	+0.00%	Tie
Specificity (Low Risk)	70.00%	70.00%	+0.00%	Tie
F1 Score (High Risk)	0.7765	0.7765	+0.0000	Tie
CV Accuracy (10-fold)	75.00%±7.78%	75.00%±7.26%	+0.00%	Tie
CV Std Dev (stability)	7.78%	7.26%	+0.52%	LR ✓

Why LDA is more appropriate than Logistic Regression for this dataset:

- **1. Performance is equal — but LDA achieves this with stronger theoretical justification.** Both methods produce identical accuracy (76.25%), sensitivity (82.50%), specificity (70.00%), and F1 score. LDA has a marginally higher AUC (0.8350 vs 0.8337). When two methods perform equally well on the same data, the one whose assumptions are better satisfied is preferred — which is LDA here.
- **2. LDA assumptions are largely met.** All 7 predictors are continuous clinical measurements (the assumption LDA needs most). 6 of 7 Levene tests pass (equal group variances). Major predictors (Age, BP, Cholesterol) are normally distributed. When these assumptions hold, LDA is statistically *more efficient* than logistic regression — it extracts more information from the covariance structure between predictors.
- **3. Perfect 50/50 class balance.** LDA is most efficient at equal class sizes. Logistic regression's advantage over LDA is handling severe class imbalance (e.g., 95%/5%). With 160 High Risk and 160 Low Risk patients, LDA is operating at its optimal condition.
- **4. Richer output: the discriminant function and D-score.** LDA produces a continuous D-score for every patient — a signed distance from the decision boundary. A patient with $D = -2.5$ is at far greater risk than $D = -0.1$, even though both are classified "High Risk." This gradient is clinically useful for triage and intervention prioritization. Logistic regression produces log-odds, which are less intuitive to rank.
- **5. Directly interpretable standardized coefficients.** LDA's standardized coefficients rank variables by importance on a common scale. Exercise is the #1 discriminator ($|coef| = 0.957$), Stress is #2 (0.932) — this is immediately clinically actionable. Logistic regression's odds ratios require additional interpretation steps and are sensitive to the scale of each predictor.
- **When logistic regression would be preferred (none apply here):** severely imbalanced classes; heavily non-normal or categorical predictors; when calibrated probability scores are the primary output needed.

11. Figures



12. Conclusions

- **Research question answered — YES.** LDA successfully discriminates between High Risk and Low Risk heart disease patients (Accuracy = 76.25%, AUC = 0.8350, CV = 75.00% $\pm$ 7.78%) — substantially above chance (50%). The model generalizes consistently across all 10 cross-validation folds.
- **The discriminant function combines all 7 variables** into a single equation (Formula B above) that can classify any new patient by plugging in their clinical measurements. The worked example ($D = -2.05 \rightarrow$ HIGH RISK) demonstrates this directly.
- **Exercise is the #1 discriminator** ($|\text{coef}| = 0.957$, structure $r = +0.473$): Low exercise is the single strongest marker of high cardiac risk. High Risk patients exercise 0.84 fewer hours per week despite this being one of the most modifiable risk factors.
- **Stress Level is #2** ($|\text{coef}| = 0.932$, structure $r = -0.449$): High Risk patients report 0.88 higher stress scores. Stress management is consistently underutilized in cardiovascular risk reduction programs.

- **Blood Pressure (#3) and Cholesterol (#4)** have nearly identical discriminating power ($|\text{coef}|$ 0.857 and 0.853), confirming their role as classic cardiovascular risk markers. BMI (#5) and Age (#6) also contribute significantly.
- **All 7 predictors are significant** in the structure matrix (all $p < 0.01$), confirming each variable independently contributes to group separation.
- **LDA and logistic regression perform identically on this dataset.** LDA is preferred because its assumptions are better satisfied (continuous predictors, near-normal distributions, balanced classes) and it produces richer output — the discriminant function with standardized coefficients and patient-level D-scores — without any performance cost.